



Faster R-CNN

R50-FPN




YOLOv10

S / X




EfficientNet

B0




DenseNet

121



Deployment engine *INT8*; in *INT8* (FP32 fallback) the FP16 parts run in FP32 with identical scales.

explicit QDQ pattern



TensorRT fuses Q/DQ into an INT8 kernel with INT32 accumulation